

# simple\_CNN\_GRAD\_CAM-128x

May 20, 2024

## 1 Analyse der Klassifizierung: 128 Pixel

```
[1]: import os
import pandas as pd
import numpy as np
import seaborn as sns

import tensorflow as tf
from tensorflow.keras import layers, models
from tensorflow.keras.callbacks import ModelCheckpoint
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.preprocessing.image import img_to_array, load_img

from sklearn.metrics import confusion_matrix, classification_report
from sklearn.model_selection import train_test_split

import matplotlib.pyplot as plt
import matplotlib.cm as cm

import cv2

#Eingabeparameter (Auflösung und Batch-size)
w, h = 128, 128
batch_size = 64

# Normalisieren der Bilder
train_data_gen = ImageDataGenerator(rescale = 1.0/255)
test_data_gen = ImageDataGenerator(rescale = 1.0/255)

# Erstellen der train- und test-Datensätze
train_Dataset = train_data_gen.flow_from_directory(
    "../images/car/preprocessed_128x/train",
    target_size = (w, h),
    batch_size = batch_size,
    class_mode = 'binary',
    shuffle = True
)
```

```

test_Dataset = test_data_gen.flow_from_directory(
    "../images/car/preprocessed_128x/test",
    target_size = (w, h),
    batch_size = batch_size,
    class_mode = 'binary',
    shuffle = False
)

# Implementierung der Modellarchitektur:
model = models.Sequential()
model.add(layers.Conv2D(16, (3, 3), activation='relu', input_shape=(w, h, 3),
    ↪name='conv2d_1'))
model.add(layers.MaxPooling2D((2, 2), name='max_pooling2d_1'))
model.add(layers.Conv2D(32, (3, 3), activation='relu', name='conv2d_2'))
model.add(layers.MaxPooling2D((2, 2), name='max_pooling2d_2'))
model.add(layers.Conv2D(64, (3, 3), activation='relu', name='conv2d_3'))
model.add(layers.MaxPooling2D((2, 2), name='max_pooling2d_3'))
model.add(layers.Flatten(name='flatten'))
model.add(layers.Dense(128, activation='relu', name='dense_1'))
model.add(layers.Dropout(0.5, name='dropout_1'))
model.add(layers.Dense(1, activation='sigmoid', name='output'))

model.compile(optimizer='adam', loss='binary_crossentropy',
    ↪metrics=['accuracy'])
model.summary()

# Speichern der besten Parameter
weight_path = './weights/'
model_checkpoint = ModelCheckpoint(
    filepath = weight_path,
    save_best_only = True,
    save_weights_only = True,
    mode = 'max',
    monitor = 'val_accuracy'
)

# Training Modell
history = model.fit(train_Dataset,
                    epochs=1,
                    validation_data=test_Dataset,
                    callbacks=[model_checkpoint]
                    )

# Laden der besten Gewichte
model.load_weights(weight_path)
# Ausführen auf dem Testdatensatz

```

```

predictions = model.predict(test_Dataset)
# Runden des Ergebnisses
binary_predictions = (predictions > 0.5).astype(int)

# Ausgabe der Bewertungsmetriken
true_labels = test_Dataset.classes
sns.heatmap(confusion_matrix(true_labels, binary_predictions), annot=True)
print(classification_report(true_labels, binary_predictions))

plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual Classes')
plt.show()

```

Found 4486 images belonging to 2 classes.


Found 1122 images belonging to 2 classes.

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 126, 126, 16)	448
max_pooling2d_1 (MaxPooling2D)	(None, 63, 63, 16)	0
conv2d_2 (Conv2D)	(None, 61, 61, 32)	4640
max_pooling2d_2 (MaxPooling2D)	(None, 30, 30, 32)	0
conv2d_3 (Conv2D)	(None, 28, 28, 64)	18496
max_pooling2d_3 (MaxPooling2D)	(None, 14, 14, 64)	0
flatten (Flatten)	(None, 12544)	0
dense_1 (Dense)	(None, 128)	1605760
dropout_1 (Dropout)	(None, 128)	0
output (Dense)	(None, 1)	129
Total params: 1629473 (6.22 MB)		
Trainable params: 1629473 (6.22 MB)		
Non-trainable params: 0 (0.00 Byte)		

71/71 [=====] - 28s 380ms/step - loss: 0.5013 -  
 accuracy: 0.7570 - val\_loss: 0.2195 - val\_accuracy: 0.9269  
 18/18 [=====] - 1s 41ms/step

	precision	recall	f1-score	support
0	0.90	0.96	0.93	561
1	0.95	0.90	0.92	561
accuracy			0.93	1122
macro avg	0.93	0.93	0.93	1122
weighted avg	0.93	0.93	0.93	1122



Folgende Labels wurden definiert:

```
[2]: train_Dataset.class_indices
```

```
[2]: {'fake': 0, 'real': 1}
```

```
[3]: def make_gradcam_heatmap(img_array, model, last_conv_layer_name,
    ↪pred_index=None):
    # Erstellen eines Modells, welches die Aktivierung der gewählten
    ↪Conv-Schicht
    # und die Ausgabe des Modells abbildet.
    grad_model = tf.keras.models.Model(
        [model.inputs], [model.get_layer(last_conv_layer_name).output, model.
    ↪output]
    )

    # Berechnung des Gradienten der am höchsten vorhergesagten Klasse für das
    ↪Eingabebild
    # in Bezug auf die Aktivierungen der letzten Conv-Schicht
    with tf.GradientTape() as tape:
        last_conv_layer_output, preds = grad_model(img_array)
        if pred_index is None:
            pred_index = tf.argmax(preds[0])
            class_channel = preds[:, pred_index]

    # Dies ist der Gradient des Ausgabeneurons
    # in Bezug auf die Ausgabefeaturemap der Conv-Schicht
    grads = tape.gradient(class_channel, last_conv_layer_output)

    # Dies ist ein Vektor, wobei jeder Eintrag die mittlere Intensität des
    ↪Gradienten
    # über einem spezifischen Featuremap-Kanal darstellt
    pooled_grads = tf.reduce_mean(grads, axis=(0, 1, 2))

    # Multiplikation jedes Kanals in der Featuremap
    # mit dem Durchschnittswert des Gradienten für diesen Kanal, um zu erfahren,
    # wie wichtig dieser Kanal für die am höchsten vorhergesagte Klasse ist.
    # Dann summieren wir alle Kanäle, um die Klassenaktivierung für die Heatmap
    ↪zu erhalten.
    last_conv_layer_output = last_conv_layer_output[0]
    heatmap = last_conv_layer_output @ pooled_grads[..., tf.newaxis]
    heatmap = tf.squeeze(heatmap)

    # Die Heatmap normalisieren
    heatmap = tf.maximum(heatmap, 0) / tf.math.reduce_max(heatmap)

    # Eine Kanal-Dimension hinzufügen und die Größe anpassen
    heatmap = tf.expand_dims(heatmap, -1) # Eine Kanal-Dimension hinzufügen
    heatmap = tf.image.resize(heatmap, (w, h)) # Heatmap in der Größe an das
    ↪Eingabebild anpassen
    heatmap = tf.squeeze(heatmap) # Die letzte Dimension entfernen, um sie
    ↪wieder 2D zu machen
```

```
return heatmap.numpy(), preds.numpy()
```

```
[8]: counter = 0

for images, labels in test_Dataset:

    for i in range(len(images)):
        if binary_predictions[i] == 0 and labels[i] == 0:
            counter = counter + 1
            image_processed = np.expand_dims(images[i], axis=0)

            plt.figure(figsize=(12, 6)) # Größere Figure, um alle Heatmaps
            ↪ sichtbar zu machen

            # Ursprüngliches Bild anzeigen
            plt.subplot(1, 4, 1)
            plt.imshow((images[i] * 255).astype("uint8"))
            plt.title(f'Original\nPredicted: {binary_predictions[i]}, Actual:
            ↪ {labels[i]}')
            plt.axis('off')

            # Heatmaps für jede Layer generieren und anzeigen
            layer_names = ['conv2d_1', 'conv2d_2', 'conv2d_3']
            for j, layer_name in enumerate(layer_names, start=2): # Startet
            ↪ bei 2, da das Originalbild auf 1 ist
                heatmap, _ = make_gradcam_heatmap(image_processed, model,
            ↪ layer_name)

                plt.subplot(1, 4, j)
                plt.imshow((images[i] * 255).astype("uint8"), alpha=0.6)
                plt.imshow(heatmap, cmap='magma', alpha=0.5) # Ensure heatmap
            ↪ is 2D

                plt.title(f'Grad-CAM\nLayer: {layer_name}')
                plt.axis('off')

            plt.show()

            if i >= 10:
                break
            if counter >= 10:
                break
```

Original  
Predicted: [0], Actual: 0.0



Grad-CAM  
Layer: conv2d\_1



Grad-CAM  
Layer: conv2d\_2



Grad-CAM  
Layer: conv2d\_3



Original  
Predicted: [0], Actual: 0.0



Grad-CAM  
Layer: conv2d\_1



Grad-CAM  
Layer: conv2d\_2




Grad-CAM  
Layer: conv2d\_3



Original  
Predicted: [0], Actual: 0.0



Grad-CAM  
Layer: conv2d\_1



Grad-CAM  
Layer: conv2d\_2



Grad-CAM  
Layer: conv2d\_3



Original  
Predicted: [0], Actual: 0.0



Grad-CAM  
Layer: conv2d\_1



Grad-CAM  
Layer: conv2d\_2




Grad-CAM  
Layer: conv2d\_3



Original  
Predicted: [0], Actual: 0.0



Grad-CAM  
Layer: conv2d\_1



Grad-CAM  
Layer: conv2d\_2



Grad-CAM  
Layer: conv2d\_3




Original  
Predicted: [0], Actual: 0.0




Grad-CAM  
Layer: conv2d\_1



Grad-CAM  
Layer: conv2d\_2



Grad-CAM  
Layer: conv2d\_3



Original  
Predicted: [0], Actual: 0.0



Grad-CAM  
Layer: conv2d\_1



Grad-CAM  
Layer: conv2d\_2



Grad-CAM  
Layer: conv2d\_3



Original  
Predicted: [0], Actual: 0.0




Grad-CAM  
Layer: conv2d\_1



Grad-CAM  
Layer: conv2d\_2



Grad-CAM  
Layer: conv2d\_3







Original  
Predicted: [0], Actual: 0.0




Grad-CAM  
Layer: conv2d\_1



Grad-CAM  
Layer: conv2d\_2



Grad-CAM  
Layer: conv2d\_3



Original  
Predicted: [0], Actual: 0.0



Grad-CAM  
Layer: conv2d\_1



Grad-CAM  
Layer: conv2d\_2



Grad-CAM  
Layer: conv2d\_3



Original  
Predicted: [0], Actual: 0.0



Grad-CAM  
Layer: conv2d\_1



Grad-CAM  
Layer: conv2d\_2



Grad-CAM  
Layer: conv2d\_3

