

simple_CNN_GRAD_CAM-256x

May 20, 2024

```
[1]: import os
import pandas as pd
import numpy as np
import seaborn as sns

import tensorflow as tf
from tensorflow.keras import layers, models
from tensorflow.keras.callbacks import ModelCheckpoint
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.preprocessing.image import img_to_array, load_img

from sklearn.metrics import confusion_matrix, classification_report
from sklearn.model_selection import train_test_split

import matplotlib.pyplot as plt
import matplotlib.cm as cm

import cv2

# Update image size to 128x128
w, h = 256, 256
batch_size = 5

train_data_gen = ImageDataGenerator(rescale = 1.0/255)
test_data_gen = ImageDataGenerator(rescale = 1.0/255)

train_Dataset = train_data_gen.flow_from_directory(
    "../images/car/preprocessed_256x/train",
    target_size = (w, h),
    batch_size = batch_size,
    class_mode = 'binary',
    shuffle = True
)

test_Dataset = test_data_gen.flow_from_directory(
    "../images/car/preprocessed_256x/test",
    target_size = (w, h),
```

```

        batch_size = batch_size,
        class_mode = 'binary',
        shuffle = False
    )

    # Update model architecture
    model = models.Sequential()
    model.add(layers.Conv2D(16, (3, 3), activation='relu', input_shape=(w, h, 3),
        ↪name='conv2d_1'))
    model.add(layers.MaxPooling2D((2, 2), name='max_pooling2d_1'))
    model.add(layers.Conv2D(32, (3, 3), activation='relu', name='conv2d_2'))
    model.add(layers.MaxPooling2D((2, 2), name='max_pooling2d_2'))
    model.add(layers.Conv2D(64, (3, 3), activation='relu', name='conv2d_3'))
    model.add(layers.MaxPooling2D((2, 2), name='max_pooling2d_3'))
    model.add(layers.Flatten(name='flatten'))
    model.add(layers.Dense(128, activation='relu', name='dense_1'))
    model.add(layers.Dropout(0.5, name='dropout_1'))
    model.add(layers.Dense(1, activation='sigmoid', name='output'))

    model.compile(optimizer='adam', loss='binary_crossentropy',
        ↪metrics=['accuracy'])
    model.summary()

    weight_path = './weights/'
    model_checkpoint = ModelCheckpoint(
        filepath = weight_path,
        save_best_only = True,
        save_weights_only = True,
        mode = 'max',
        monitor = 'val_accuracy'
    )
    history = model.fit(train_Dataset,
                        epochs=2,
                        validation_data=test_Dataset,
                        callbacks=[model_checkpoint]
    )

    # Load best weights and make predictions
    model.load_weights(weight_path)
    predictions = model.predict(test_Dataset)
    binary_predictions = (predictions > 0.5).astype(int)

    # Generate confusion matrix and classification report
    true_labels = test_Dataset.classes
    sns.heatmap(confusion_matrix(true_labels, binary_predictions), annot=True)
    print(classification_report(true_labels, binary_predictions))

```

```
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual Classes')
plt.show()
```

Found 4486 images belonging to 2 classes.

Found 1122 images belonging to 2 classes.

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 254, 254, 16)	448
max_pooling2d_1 (MaxPooling2D)	(None, 127, 127, 16)	0
conv2d_2 (Conv2D)	(None, 125, 125, 32)	4640
max_pooling2d_2 (MaxPooling2D)	(None, 62, 62, 32)	0
conv2d_3 (Conv2D)	(None, 60, 60, 64)	18496
max_pooling2d_3 (MaxPooling2D)	(None, 30, 30, 64)	0
flatten (Flatten)	(None, 57600)	0
dense_1 (Dense)	(None, 128)	7372928
dropout_1 (Dropout)	(None, 128)	0
output (Dense)	(None, 1)	129

Total params: 7396641 (28.22 MB)

Trainable params: 7396641 (28.22 MB)

Non-trainable params: 0 (0.00 Byte)

Epoch 1/2

898/898 [=====] - 127s 141ms/step - loss: 0.3302 - accuracy: 0.8535 - val_loss: 0.2799 - val_accuracy: 0.8939


Epoch 2/2

898/898 [=====] - 116s 129ms/step - loss: 0.1734 - accuracy: 0.9360 - val_loss: 0.1161 - val_accuracy: 0.9563

225/225 [=====] - 4s 18ms/step

precision recall f1-score support

0	0.95	0.96	0.96	561
1	0.96	0.95	0.96	561
accuracy			0.96	1122
macro avg	0.96	0.96	0.96	1122
weighted avg	0.96	0.96	0.96	1122



```
[2]: train_Dataset.class_indices
```

```
[2]: {'fake': 0, 'real': 1}
```

```
[3]: def make_gradcam_heatmap(img_array, model, last_conv_layer_name,
    ↪pred_index=None):
    # Erstellen eines Modells, welches die Aktivierung der gewählten
    ↪Conv-Schicht
    # und die Ausgabe des Modells abbildet.
    grad_model = tf.keras.models.Model(
        [model.inputs], [model.get_layer(last_conv_layer_name).output, model.
    ↪output])
```

```

)

# Berechnung des Gradienten der am höchsten vorhergesagten Klasse für das
↳Eingabebild
# in Bezug auf die Aktivierungen der letzten Conv-Schicht
with tf.GradientTape() as tape:
    last_conv_layer_output, preds = grad_model(img_array)
    if pred_index is None:
        pred_index = tf.argmax(preds[0])
    class_channel = preds[:, pred_index]

# Dies ist der Gradient des Ausgabeneurons
# in Bezug auf die Ausgabefeaturemap der Conv-Schicht
grads = tape.gradient(class_channel, last_conv_layer_output)

# Dies ist ein Vektor, wobei jeder Eintrag die mittlere Intensität des
↳Gradienten
# über einem spezifischen Featuremap-Kanal darstellt
pooled_grads = tf.reduce_mean(grads, axis=(0, 1, 2))

# Multiplikation jedes Kanals in der Featuremap
# mit dem Durchschnittswert des Gradienten für diesen Kanal, um zu erfahren,
# wie wichtig dieser Kanal für die am höchsten vorhergesagte Klasse ist.
# Dann summieren wir alle Kanäle, um die Klassenaktivierung für die Heatmap
↳zu erhalten.
last_conv_layer_output = last_conv_layer_output[0]
heatmap = last_conv_layer_output @ pooled_grads[..., tf.newaxis]
heatmap = tf.squeeze(heatmap)

# Die Heatmap normalisieren
heatmap = tf.maximum(heatmap, 0) / tf.math.reduce_max(heatmap)

# Eine Kanal-Dimension hinzufügen und die Größe anpassen
heatmap = tf.expand_dims(heatmap, -1) # Eine Kanal-Dimension hinzufügen
heatmap = tf.image.resize(heatmap, (w, h)) # Heatmap in der Größe an das
↳Eingabebild anpassen
heatmap = tf.squeeze(heatmap) # Die letzte Dimension entfernen, um sie
↳wieder 2D zu machen

return heatmap.numpy(), preds.numpy()

```

```

[6]: counter = 0

for images, labels in test_Dataset:

    for i in range(len(images)):
        if binary_predictions[i] == 0 and labels[i] == 0:

```

```

counter = counter + 1
image_processed = np.expand_dims(images[i], axis=0)

plt.figure(figsize=(12, 6)) # Größere Figure, um alle Heatmaps
↪ sichtbar zu machen

# Ursprüngliches Bild anzeigen
plt.subplot(1, 4, 1)
plt.imshow((images[i] * 255).astype("uint8"))
plt.title(f'Original\nPredicted: {binary_predictions[i]}, Actual:
↪ {labels[i]}')
plt.axis('off')

# Heatmaps für jede Layer generieren und anzeigen
layer_names = ['conv2d_1', 'conv2d_2', 'conv2d_3']
for j, layer_name in enumerate(layer_names, start=2): # Startet
↪ bei 2, da das Originalbild auf 1 ist
    heatmap, _ = make_gradcam_heatmap(image_processed, model,
↪ layer_name)

    plt.subplot(1, 4, j)
    plt.imshow((images[i] * 255).astype("uint8"), alpha=0.6)
    plt.imshow(heatmap, cmap='magma', alpha=0.5) # Ensure heatmap
↪ is 2D

    plt.title(f'Grad-CAM\nLayer: {layer_name}')
    plt.axis('off')

plt.show()

if i >= 10:
    break
if counter >= 10:
    break

```

Original
Predicted: [0], Actual: 0.0



Grad-CAM
Layer: conv2d_1



Grad-CAM
Layer: conv2d_2



Grad-CAM
Layer: conv2d_3



Original
Predicted: [0], Actual: 0.0



Grad-CAM
Layer: conv2d_1



Grad-CAM
Layer: conv2d_2



Grad-CAM
Layer: conv2d_3



Original
Predicted: [0], Actual: 0.0



Grad-CAM
Layer: conv2d_1



Grad-CAM
Layer: conv2d_2




Grad-CAM
Layer: conv2d_3




Original
Predicted: [0], Actual: 0.0



Grad-CAM
Layer: conv2d_1



Grad-CAM
Layer: conv2d_2



Grad-CAM
Layer: conv2d_3



Original
Predicted: [0], Actual: 0.0



Grad-CAM
Layer: conv2d_1



Grad-CAM
Layer: conv2d_2



Grad-CAM
Layer: conv2d_3



Original
Predicted: [0], Actual: 0.0



Grad-CAM
Layer: conv2d_1




Grad-CAM
Layer: conv2d_2




Grad-CAM
Layer: conv2d_3



Original
Predicted: [0], Actual: 0.0



Grad-CAM
Layer: conv2d_1




Grad-CAM
Layer: conv2d_2




Grad-CAM
Layer: conv2d_3



Original
Predicted: [0], Actual: 0.0




Grad-CAM
Layer: conv2d_1



Grad-CAM
Layer: conv2d_2



Grad-CAM
Layer: conv2d_3



Original
Predicted: [0], Actual: 0.0



Grad-CAM
Layer: conv2d_1



Grad-CAM
Layer: conv2d_2




Grad-CAM
Layer: conv2d_3



Original
Predicted: [0], Actual: 0.0



Grad-CAM
Layer: conv2d_1



Grad-CAM
Layer: conv2d_2



Grad-CAM
Layer: conv2d_3



Original
Predicted: [0], Actual: 0.0



Grad-CAM
Layer: conv2d_1



Grad-CAM
Layer: conv2d_2



Grad-CAM
Layer: conv2d_3



Original
Predicted: [0], Actual: 0.0



Grad-CAM
Layer: conv2d_1



Grad-CAM
Layer: conv2d_2



Grad-CAM
Layer: conv2d_3

