

simple_CNN_GRAD_CAM-512x1

May 20, 2024

```
[1]: import os
import pandas as pd
import numpy as np
import seaborn as sns

import tensorflow as tf
from tensorflow.keras import layers, models
from tensorflow.keras.callbacks import ModelCheckpoint
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.preprocessing.image import img_to_array, load_img

from sklearn.metrics import confusion_matrix, classification_report
from sklearn.model_selection import train_test_split

import matplotlib.pyplot as plt
import matplotlib.cm as cm

import cv2

# Update image size to 128x128
w, h = 512, 512
batch_size = 2

train_data_gen = ImageDataGenerator(rescale = 1.0/255)
test_data_gen = ImageDataGenerator(rescale = 1.0/255)

train_Dataset = train_data_gen.flow_from_directory(
    "../images/car/preprocessed_512x/train",
    target_size = (w, h),
    batch_size = batch_size,
    class_mode = 'binary',
    shuffle = True
)

test_Dataset = test_data_gen.flow_from_directory(
    "../images/car/preprocessed_512x/test",
    target_size = (w, h),
```

```

        batch_size = batch_size,
        class_mode = 'binary',
        shuffle = False
    )

    # Update model architecture
    model = models.Sequential()
    model.add(layers.Conv2D(16, (3, 3), activation='relu', input_shape=(w, h, 3),
        ↪name='conv2d_1'))
    model.add(layers.MaxPooling2D((2, 2), name='max_pooling2d_1'))
    model.add(layers.Conv2D(32, (3, 3), activation='relu', name='conv2d_2'))
    model.add(layers.MaxPooling2D((2, 2), name='max_pooling2d_2'))
    model.add(layers.Conv2D(64, (3, 3), activation='relu', name='conv2d_3'))
    model.add(layers.MaxPooling2D((2, 2), name='max_pooling2d_3'))
    model.add(layers.Flatten(name='flatten'))
    model.add(layers.Dense(128, activation='relu', name='dense_1'))
    model.add(layers.Dropout(0.5, name='dropout_1'))
    model.add(layers.Dense(1, activation='sigmoid', name='output'))

    model.compile(optimizer='adam', loss='binary_crossentropy',
        ↪metrics=['accuracy'])
    model.summary()

    weight_path = './weights/'
    model_checkpoint = ModelCheckpoint(
        filepath = weight_path,
        save_best_only = True,
        save_weights_only = True,
        mode = 'max',
        monitor = 'val_accuracy'
    )
    history = model.fit(train_Dataset,
                        epochs=1,
                        validation_data=test_Dataset,
                        callbacks=[model_checkpoint]
                    )

    # Load best weights and make predictions
    model.load_weights(weight_path)
    predictions = model.predict(test_Dataset)
    binary_predictions = (predictions > 0.5).astype(int)

    # Generate confusion matrix and classification report
    true_labels = test_Dataset.classes
    sns.heatmap(confusion_matrix(true_labels, binary_predictions), annot=True)
    print(classification_report(true_labels, binary_predictions))

```

```
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual Classes')
plt.show()
```

Found 4486 images belonging to 2 classes.

Found 1122 images belonging to 2 classes.

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 510, 510, 16)	448
max_pooling2d_1 (MaxPooling2D)	(None, 255, 255, 16)	0
conv2d_2 (Conv2D)	(None, 253, 253, 32)	4640
max_pooling2d_2 (MaxPooling2D)	(None, 126, 126, 32)	0
conv2d_3 (Conv2D)	(None, 124, 124, 64)	18496
max_pooling2d_3 (MaxPooling2D)	(None, 62, 62, 64)	0
flatten (Flatten)	(None, 246016)	0
dense_1 (Dense)	(None, 128)	31490176
dropout_1 (Dropout)	(None, 128)	0
output (Dense)	(None, 1)	129

Total params: 31513889 (120.22 MB)

Trainable params: 31513889 (120.22 MB)

Non-trainable params: 0 (0.00 Byte)


2243/2243 [=====] - 794s 354ms/step - loss: 0.3657 -

accuracy: 0.8694 - val_loss: 0.2273 - val_accuracy: 0.9064

561/561 [=====] - 17s 30ms/step

	precision	recall	f1-score	support
0	0.95	0.86	0.90	561
1	0.87	0.95	0.91	561
accuracy			0.91	1122

macro avg	0.91	0.91	0.91	1122
weighted avg	0.91	0.91	0.91	1122



```
[2]: train_Dataset.class_indices
```

```
[2]: {'fake': 0, 'real': 1}
```

```
[3]: def make_gradcam_heatmap(img_array, model, last_conv_layer_name,
    ↪pred_index=None):
    grad_model = tf.keras.models.Model(
        [model.inputs], [model.get_layer(last_conv_layer_name).output, model.
    ↪output]
    )

    with tf.GradientTape() as tape:
        last_conv_layer_output, preds = grad_model(img_array)
        if pred_index is None:
            pred_index = tf.argmax(preds[0])
        class_channel = preds[:, pred_index]
```

```

grads = tape.gradient(class_channel, last_conv_layer_output)

pooled_grads = tf.reduce_mean(grads, axis=(0, 1, 2))
last_conv_layer_output = last_conv_layer_output[0]
heatmap = last_conv_layer_output @ pooled_grads[..., tf.newaxis]
heatmap = tf.squeeze(heatmap)

# Normalize the heatmap
heatmap = tf.maximum(heatmap, 0) / tf.math.reduce_max(heatmap)

# Add a channel dimension and resize
heatmap = tf.expand_dims(heatmap, -1) # Add a channel dimension
heatmap = tf.image.resize(heatmap, (w, h)) # Resize heatmap to match the
↪ input image size
heatmap = tf.squeeze(heatmap) # Remove the last dimension to make it 2D
↪ again

return heatmap.numpy(), preds.numpy()

```

```

[9]: counter = 0

for images, labels in test_Dataset:

    for i in range(len(images)):
        if binary_predictions[i] == 0 and labels[i] == 0:
            counter = counter + 1
            image_processed = np.expand_dims(images[i], axis=0)

            plt.figure(figsize=(12, 6)) # Größere Figure, um alle Heatmaps
            ↪ sichtbar zu machen

            # Ursprüngliches Bild anzeigen
            plt.subplot(1, 4, 1)
            plt.imshow((images[i] * 255).astype("uint8"))
            plt.title(f'Original\nPredicted: {binary_predictions[i]}, Actual: {
            ↪ labels[i]}')
            plt.axis('off')

            # Heatmaps für jede Layer generieren und anzeigen
            layer_names = ['conv2d_1', 'conv2d_2', 'conv2d_3']
            for j, layer_name in enumerate(layer_names, start=2): # Startet
            ↪ bei 2, da das Originalbild auf 1 ist
                heatmap, _ = make_gradcam_heatmap(image_processed, model,
            ↪ layer_name)

                plt.subplot(1, 4, j)

```

```

plt.imshow((images[i] * 255).astype("uint8"), alpha=0.6)
plt.imshow(heatmap, cmap='magma', alpha=0.5) # Ensure heatmap
↳ is 2D

plt.title(f'Grad-CAM\nLayer: {layer_name}')
plt.axis('off')

plt.show()

if i >= 10:
    break
if counter >= 10:
    break

```

Original
Predicted: [0], Actual: 0.0



Grad-CAM
Layer: conv2d_1



Grad-CAM
Layer: conv2d_2



Grad-CAM
Layer: conv2d_3



Original
Predicted: [0], Actual: 0.0



Grad-CAM
Layer: conv2d_1



Grad-CAM
Layer: conv2d_2



Grad-CAM
Layer: conv2d_3



Original
Predicted: [0], Actual: 0.0



Grad-CAM
Layer: conv2d_1



Grad-CAM
Layer: conv2d_2




Grad-CAM
Layer: conv2d_3




Original
Predicted: [0], Actual: 0.0




Grad-CAM
Layer: conv2d_1



Grad-CAM
Layer: conv2d_2



Grad-CAM
Layer: conv2d_3



Original
Predicted: [0], Actual: 0.0



Grad-CAM
Layer: conv2d_1



Grad-CAM
Layer: conv2d_2



Grad-CAM
Layer: conv2d_3



Original
Predicted: [0], Actual: 0.0




Grad-CAM
Layer: conv2d_1



Grad-CAM
Layer: conv2d_2



Grad-CAM
Layer: conv2d_3



Original
Predicted: [0], Actual: 0.0



Grad-CAM
Layer: conv2d_1



Grad-CAM
Layer: conv2d_2



Grad-CAM
Layer: conv2d_3



Original
Predicted: [0], Actual: 0.0




Grad-CAM
Layer: conv2d_1



Grad-CAM
Layer: conv2d_2




Grad-CAM
Layer: conv2d_3




Original
Predicted: [0], Actual: 0.0



Grad-CAM
Layer: conv2d_1



Grad-CAM
Layer: conv2d_2



Grad-CAM
Layer: conv2d_3



Original
Predicted: [0], Actual: 0.0



Grad-CAM
Layer: conv2d_1



Grad-CAM
Layer: conv2d_2



Grad-CAM
Layer: conv2d_3

